

[Get unlimited access](#)[Open in app](#)

Eliud Nduati [Following](#)

Feb 7 · 4 min read · [Listen](#)



Save



Week #4 in Machine Learning

Diabetes classification — Supervised ML classification problem



[source](#)

In this notebook, I am applying supervised machine learning classifications to a diabetic dataset. I aim to determine whether the tested data has diabetes or not. I will use KNN, decision tree, random forest, Support vector machine, logistic regression and Naive Bayes algorithms. I will also make evaluation of all the models used by using confusion matrix.

Data Loading

```
1 # import libraries
2 import pandas as pd
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import seaborn as sns
6
7 from sklearn.model_selection import train_test_split
8 from sklearn.linear_model import LogisticRegression # Logistic regression
```

6e350047-034b-4b02-a70a-c794f3bca571.py hosted with ❤ by GitHub

[view raw](#)

```
1 # loading the dataset
2 df = pd.read_csv("diabetes.csv")
3 df.head()
```



[Get unlimited access](#)[Open in app](#)

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

```
df.shape
```

```
(768, 9)
```

Descriptive statistics

This step shows the decriptive statistics of all the numerical columns in the dataset.

```
df.describe()
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
count	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000	768.000000
mean	3.845052	120.894531	69.105469	20.536458	79.799479	31.992578	0.471876	33.240885	0.348958
std	3.369578	31.972618	19.355807	15.952218	115.244002	7.884160	0.331329	11.760232	0.476951
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.078000	21.000000	0.000000
25%	1.000000	99.000000	62.000000	0.000000	0.000000	27.300000	0.243750	24.000000	0.000000
50%	3.000000	117.000000	72.000000	23.000000	30.500000	32.000000	0.372500	29.000000	0.000000
75%	6.000000	140.250000	80.000000	32.000000	127.250000	36.600000	0.626250	41.000000	1.000000
max	17.000000	199.000000	122.000000	99.000000	846.000000	67.100000	2.420000	81.000000	1.000000

EDA

Exploratory Data Analysis

This step focuses on exploring throuhg the data to determine the data types, check for missing data point and fix them etc.


```
1 # Smmarize the characteristics of the data coluhmns to check for data types and missing values
2 df.info()
```

66d60a71-946a-4b77-a2e4-0026fd1375f3.py hosted with ❤ by GitHub

[view raw](#)

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Pregnancies           768 non-null   int64
 1   Glucose               768 non-null   int64
 2   BloodPressure         768 non-null   int64
 3   SkinThickness         768 non-null   int64
 4   Insulin               768 non-null   int64
 5   BMI                   768 non-null   float64
 6   DiabetesPedigreeFunction 768 non-null   float64
 7   Age                   768 non-null   int64
 8   Outcome               768 non-null   int64
```



[Get unlimited access](#)[Open in app](#)

This chart shows how the Outcomes are classified. There are more Negative outcomes, 0, than positive outcomes, 1. There's not much cleaning to be done on our dataset. We can therefore go directly to the Machine learning processes.

Machine Learning

```
1 #Have the outcome data as y
2 y = df.Outcome.values
3
4 # remove the Outcome data from the dataset and have the remaining as x
5 x_data = df.drop(['Outcome'], axis=1)
```

7a6fe98c-f53d-42b4-b808-a1f7ee732ee3.py hosted with ❤ by GitHub

[view raw](#)

```
x_data.head()
```



[Get unlimited access](#)[Open in app](#)

	0	148	72	88	0	88.8	0.324	88
1	1	85	66	29	0	26.6	0.351	31
2	8	183	64	0	0	23.3	0.672	32
3	1	89	66	23	94	28.1	0.167	21
4	0	137	40	35	168	43.1	2.288	33

```
y[1:10]
array([0, 1, 0, 1, 0, 1, 0, 1, 1])
```

```
1 #Normalization to handle unbalanced features
2 x = (x_data - np.min(x_data))/(np.max(x_data) - np.min(x_data)).values
```

7056a021-5dc2-422c-b78d-ef29927702f0.py hosted with ❤ by GitHub

[view raw](#)

```
1 # Split the data into training and test set. We use 20% test data with a random state of 42
2 x_train, x_test, y_train, y_test = train_test_split(x, y, test_size= 0.2, random_state=42)
```

64d24f1b-49a9-4bbe-9c6f-f038a1a9ae67.py hosted with ❤ by GitHub

[view raw](#)

Logistic Regression Classification

Logistic regression is a powerful algorithm when you have a binary classification problem

```
1 lr = LogisticRegression()
2 lr.fit(x_train, y_train)
3 print("test accuracy {}".format(lr.score(x_test, y_test)))
4
5 lr_score=lr.score(x_test, y_test)
```

8023f7fd-4f14-4dda-b210-c6c0d4dbab96.py hosted with ❤ by GitHub


[view raw](#)

```
test accuracy 0.7662337662337663
```

```
1 # using confusion matrix to evaluate the linear regression
2 from sklearn.metrics import confusion_matrix
3
4 y_prediction = lr.predict(x_test)
5 y_actual = y_test
6 cm = confusion_matrix(y_actual, y_prediction)
7
8 # heatmap visulization of confusion matrix
9 f, ax = plt.subplots(figsize =(5, 5))
10 sns.heatmap(cm, annot = True, linewidth=1, linecolor="green", fmt =".0f", ax=ax)
11 plt.xlabel("y_prediction")
12 plt.ylabel("y_actual")
13 plt.show()
```

c983ccaf-0b36-4d50-86f0-2faab4b5af82.py hosted with ❤ by GitHub

[view raw](#)



KNN Classification

We need to choose a small k value but not too small that it causes overfitting while big k value causes underfitting. The K value we choose needs to be as close to our test points as possible. For this case, we use the standard k value which is k=3

```
1 # import KNN classification model
2 from sklearn.neighbors import KNeighborsClassifier
3 k=11
4 knn = KNeighborsClassifier(n_neighbors=k)
5 knn.fit(x_train, y_train)
6 prediction = knn.predict(x_test)
7 print("{} nn score: {}".format(k, knn.score(x_test, y_test)))
8
9 knn_score = knn.score(x_test, y_test)
```

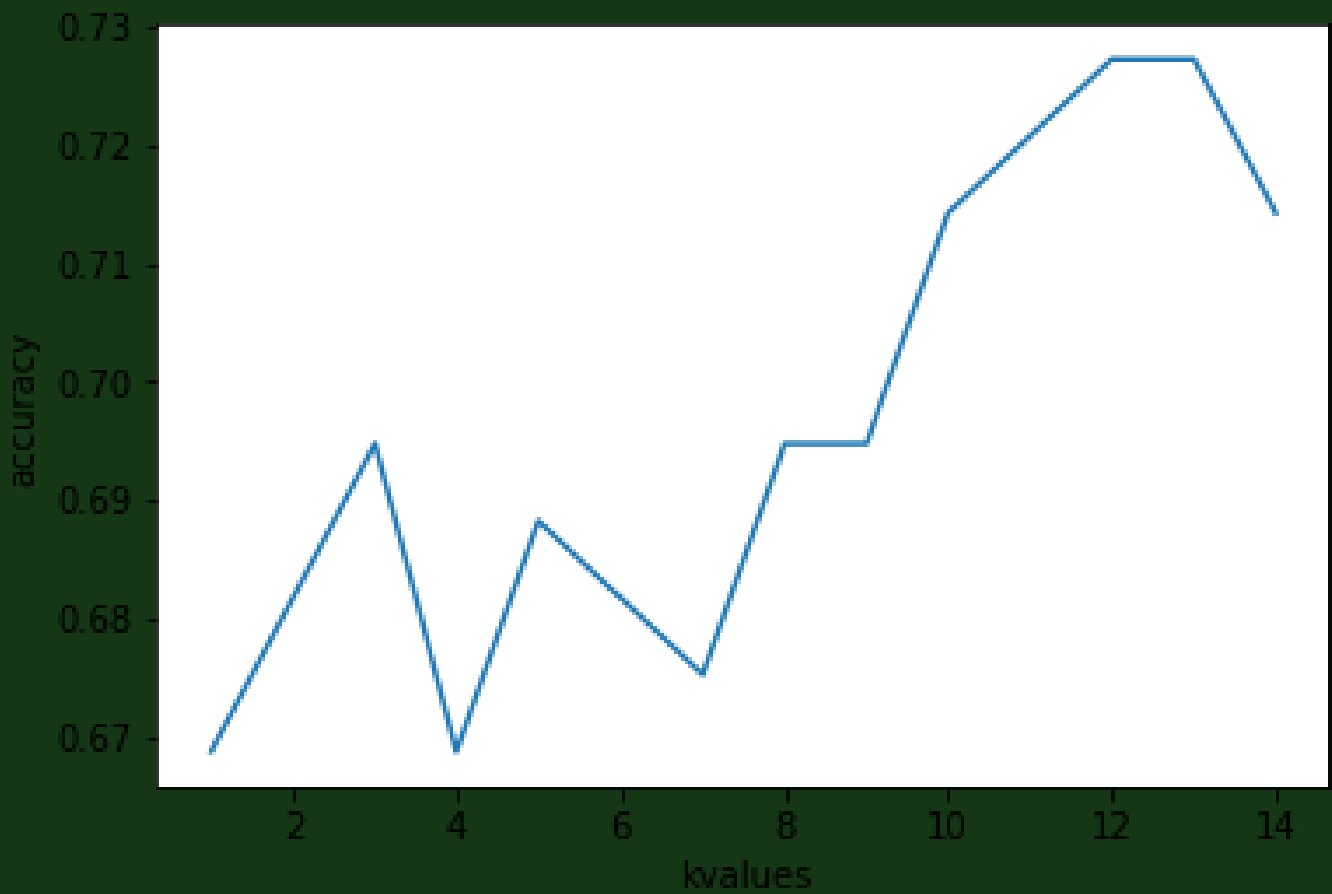


[Get unlimited access](#)[Open in app](#)

```
11 nn score: 0.7207792207792207
```

```
1 # testing different values of k with accuracy to determine the most favorable
2 # k ranges from 1-15
3 score_list = []
4 for each in range(1, 15):
5     knn2 = KNeighborsClassifier(n_neighbors= each)
6     knn2.fit(x_train, y_train)
7     score_list.append(knn2.score(x_test, y_test))
8
9 plt.plot(range(1, 15), score_list)
10 plt.xlabel("kvalues")
11 plt.ylabel("accuracy")
12 plt.show()
```

c62ef9db-2629-494c-b558-0f4cc0c11b36.py hosted with ❤ by GitHub

[view raw](#)


K=11, 12 gives the best accuracy for our case



[Get unlimited access](#)[Open in app](#)

```
6 # Heatmap visualization of confusion matrix
7 f, ax = plt.subplots(figsize = (5, 5))
8 sns.heatmap(cm, annot = True, linewidths=1, linecolor = "green", fmt = ".0f", ax=ax)
9 plt.xlabel("y_prediction")
10 plt.ylabel("y_actual")
11 plt.show()
```

6799f319-4ece-47c8-9c5a-76dd4fa80874.py hosted with ❤ by GitHub

[view raw](#)

Decision Tree Classification



[Get unlimited access](#)[Open in app](#)

```
3 dt = DecisionTreeClassifier(random_state = 42)
4 dt.fit(x_train, y_train)
5
6 print("score: ", dt.score(x_test, y_test))
7
8 dt_score = dt.score(x_test, y_test)
```

c6d00de1-2b0d-4098-a067-a147eedcc20.py hosted with ❤ by GitHub


[view raw](#)

```
score: 0.7467532467532467
```

```
1 # confsion matrix
2 y_prediction = dt.predict(x_test)
3 y_actual = y_test
4 cm = confusion_matrix(y_actual, y_prediction)
5
6 # heatmap
7 f, ax = plt.subplots(figsize = (5, 5))
8 sns.heatmap(cm, annot = True, linewidths=1, linecolor="black", fmt = ".0f", ax=ax)
9 plt.xlabel("y_prediction")
10 plt.ylabel("y_actual")
11 plt.show()
```

71219024-3810-4f4a-9e93-752c2168037d.py hosted with ❤ by GitHub

[view raw](#)

[Get unlimited access](#)[Open in app](#)

Random Forest Classification

This approach uses multiple decision trees and take the average of the results of these decision trees. this average is used to determine the class of the test points. It is one of the ensemble methods that uses multiple classes to predict the target

```
1 # random forest classification
2 from sklearn.ensemble import RandomForestClassifier
3 # set n_estimators to 100 which means the model will use 100 subsets
4 rf = RandomForestClassifier(n_estimators = 100, random_state = 42)
5 rf.fit(x_train, y_train)
6 print("random forest model score: ", rf.score(x_test, y_test))
7 rf_score= rf.score(x_test, y_test)
```


112fe9e2-9868-405a-b659-49dc8219f6ac.py hosted with ❤ by GitHub

[view raw](#)

[Get unlimited access](#)[Open in app](#)

```
3 y_actual = y_test
4 cm = confusion_matrix(y_actual, y_prediction)
5
6 # Heatmap
7 f, ax = plt.subplots(figsize=(5,5))
8 sns.heatmap(cm, annot = True, linewidths=0.5, linecolor = "green", fmt = ".0f", ax=ax)
9 plt.xlabel("y_prediction")
10 plt.ylabel("y_actual")
11 plt.show()
```

61f21c5f-5246-4e1a-bb8b-f6e982d7a7f2.py hosted with ❤ by GitHub

[view raw](#)

[Get unlimited access](#)[Open in app](#)

best value

```
1 # SVM approach
2 from sklearn.svm import SVC
3 svm = SVC(random_state = 42)
4 svm.fit(x_train, y_train)
5
6 print("Accuracy of SVM: ", svm.score(x_test, y_test))
7
8 svm_score = svm.score(x_test, y_test)
```

678e3eb7-8f3d-430e-afc4-37fd870f29ca.py hosted with ❤ by GitHub

[view raw](#)

Accuracy of SVM: 0.7467532467532467


```
1 # Confusion matrix
2 y_prediction = svm.predict(x_test)
3 y_actual = y_test
4 cm = confusion_matrix(y_actual, y_prediction)
5
6 # Heatmap
7 f, ax = plt.subplots(figsize=(5,5))
8 sns.heatmap(cm, annot = True, linewidths=0.5, linecolor = "green", fmt = ".0f", ax=ax)
9 plt.xlabel("y_prediction")
10 plt.ylabel("y_actual")
11 plt.show()
```

4c6aae72-3cd2-4a8e-ac85-67e4980caed8.py hosted with ❤ by GitHub

[view raw](#)

42



[Get unlimited access](#)[Open in app](#)

Naive Bayes Classification

this is a probabilistic classifier which applies Bayes theorem with strong independence assumption between the features. It works by determining similarity range and calculating probability of the X points in the A feature $P(A_feature|x)$

```
1 # Naive Bayes
2 from sklearn.naive_bayes import GaussianNB
3 nb = GaussianNB()
4 nb.fit(x_train, y_train)
5
6 print("accuracy of naive bayes: ", nb.score(x_test, y_test))
7
8 nb_score = nb.score(x_test, y_test)
```

b3defc57-ccc9-42d9-bceb-a5aee9ba8fcd.py hosted with ❤ by GitHub


[view raw](#)

[Get unlimited access](#)[Open in app](#)

```
1 # Confusion matrix
2 y_prediction = nb.predict(x_test)
3 y_actual = y_test
4 cm = confusion_matrix(y_actual, y_prediction)
5
6 # heatmap
7 f, ax = plt.subplots(figsize=(5,5))
8 sns.heatmap(cm, annot = True, linewidths=0.5, linecolor = "green", fmt = ".0f", ax=ax)
9 plt.xlabel("y_prediction")
10 plt.ylabel("y_actual")
11 plt.show()
```

d49ccaa0-5204-4509-853e-e0bf39128db3.py hosted with ❤ by GitHub

[view raw](#)

[Get unlimited access](#)[Open in app](#)


Comparision Using Confusion matrix

Below I visualize all confusion matrices to all classifiers

```
1 class_name = ("Logistic Regression", "KNN", "Decision Tree", "Random Forest", "SVM", "Naive bayes")
2 class_score = (lr_score, knn_score, dt_score, rf_score, svm_score, nb_score)
3 y_pos = np.arange(len(class_score))
4 colors = ("red", "green", "gold", "blue", "orange", "black")
5 plt.figure(figsize=(18, 12))
6 plt.bar(y_pos, class_score, color=colors)
7 plt.xticks(y_pos, class_name, fontsize=20)
8 plt.yticks(np.arange(0.00, 1.05, step = 0.05))
9 plt.ylabel("Accuracy")
10
11 plt.title("Confusion Matrix Comparision of Classifiers", fontsize=14)
12 plt.savefig("graph.png")
13 plt.show()
```

3d0d7c72-5e1e-4089-aba6-937d68bbbed69.py hosted with ❤ by GitHub

[view raw](#)

[Get unlimited access](#)[Open in app](#)

Conclusion

From the past 3 week's posts, we have looked at the theoretical part of classification algorithms, here we have applied the algorithms. Check out for next weeks post.

Enjoy the read? Reward the writer. Beta

Your tip will go to Eliud Nduati through a third-party platform of their choice, letting them know you appreciate their story.

[Give a tip](#)

Get an email whenever Eliud Nduati publishes.

Hello friend! Feel free to respond on your views and concerns. Let's engage!!

Emails will be sent to adunoluwa1@gmail.com. [Not you?](#)

[Subscribe](#)



Get unlimited access

Open in app

